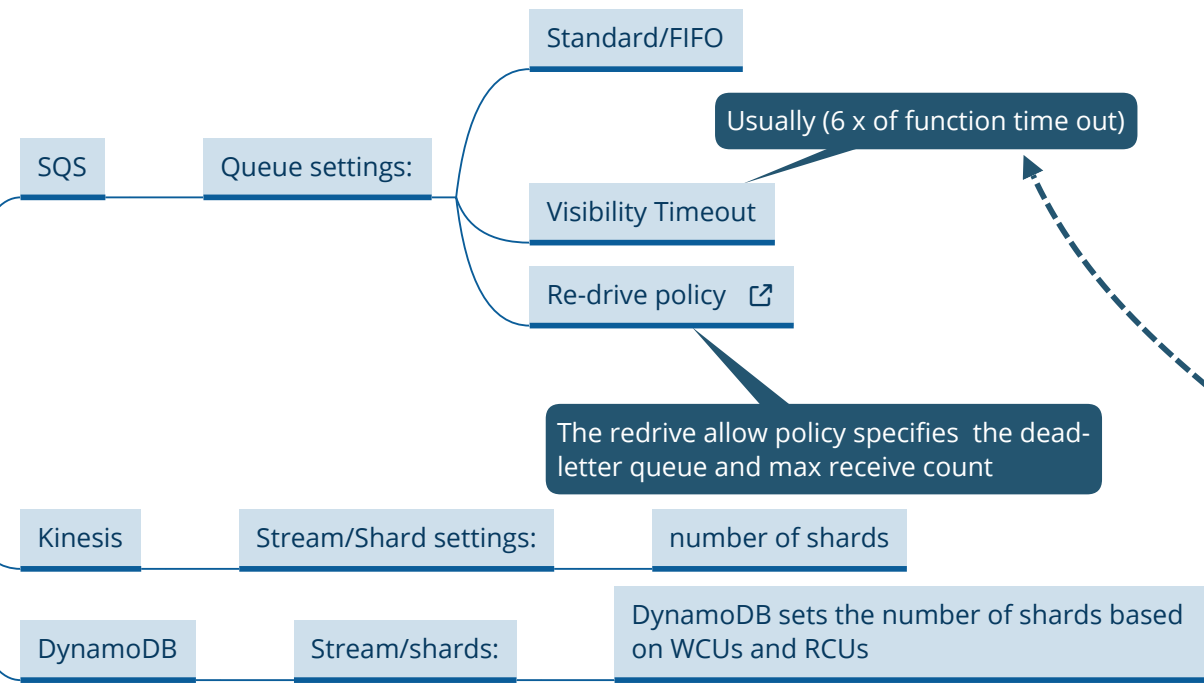


Event Polling

Even Sources



The lambda service polls the messages from queue or stream . This is "synchronous" call

Polling by Lambda service

- The event sources "push" the events into the queue or stream
- The "lambda service" periodically polls the events from the queue or stream
- Every poll gets a "batch" of events from queue or stream

We have some control in configuring the lambda that is polling queue and streams such as:

- Function timeout
- Function batch size (how many messages to pick up per batch)- from 1 to 10
 - Larger batch size reduce the cost of polling and improve efficiency for the workloads that are processed fast
 - For longer running functions it's better we set "lower batch size" to make sure all the messages are processed before function times out
- Design for Idempotency
- Code for partial failure
- Batch Window (max 300s)
 - Increasing Batch Window decreases the the number of polls

Not setting the correct batch size causes unsuccessful batch processing

Error Handling between Lambda service and event sources

- The lambda service has a built-in mechanism to retry and to get the events from the queue or stream
- See the "Lambda" tab for more info

Events in queue

- Each message is independent
- Rate of getting the messages are different, sometimes we have many sometime none
- Each message is for one consumer
 - Lambda by default starts polling in batches.
 - "5 polling processes"
 - the size of each batch to poll (the # of messages in each batch) is configurable
- Each message in the batch has to be deleted after processing
 - Lambda service does that for you (you do not need to write code for that) if the batch is processes successfully
- Error Handling
 - See the "Lambda" Tab, Error Handling , Polling (event) mode
 - If processing one message in a batch fails, the entire batch is visible again to be retried in the next batch poll
 - If this behavior is not desired, you need to delete the the messages programmatically when the message is a batch is processes
 - Also see how setting a wrong batch size can cause problem in batch processing
- Scaling
 - If the messages queue gets larger, the Lambda adjusts its number of polling processes
 - If there are many error in processing the messages, Lambda scales down, assuming that there is a problem in downstream consumers

Events in streams

- Some messages are related together and they belong to one core entity
 - So we need to process them together
- Usually we have a high rate of volume of messages
- Each message can be consumed by more than one consumer
 - We "do not" delete a message after processing as another consumer may need it again
- Error Handling
 - See the Lambda tab, "Error Handling, Polling (event) mode branch
 - Lambda will not increase the concurrent invocations unless you increase the number of shards
- Scaling/Performance
 - We can use concurrent batches per shard
 - This is called "Parallelization factor"
 - Batch window
 - By Default: Lambda reads records from a stream at a fixed cadence (i.e., once per second for Kinesis data streams) and invokes a function with a batch of records
 - Batch Window setting in stream allows you to wait as long as 300s to build a batch before invoking a function
 - Polling Rate
 - Lambda polls each shard 4 times per second.

Poll model (Even based).

